

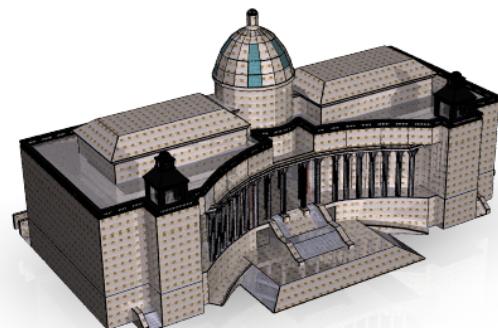
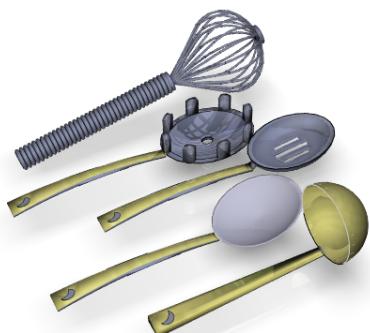
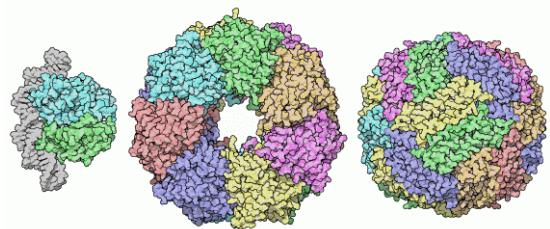
CSE 152: Computer Vision

Hao Su

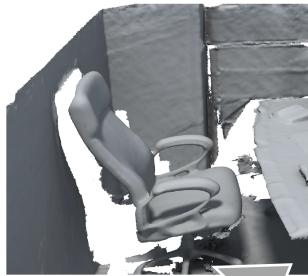
Lecture 12: 3D Deep Learning



Credit: Stanford CS231n, L13



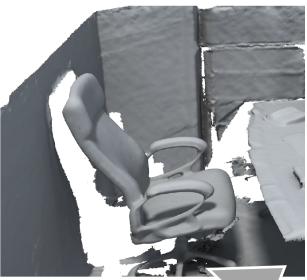
Broad Applications of 3D data



Robotics



Broad Applications of 3D data



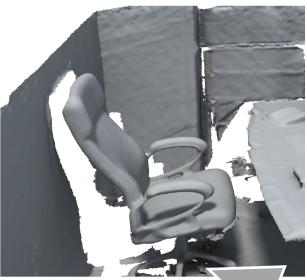
Robotics



**Augmented
Reality**



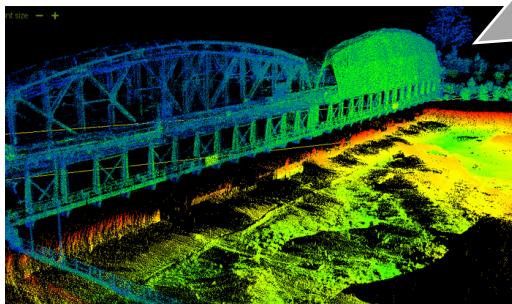
Broad Applications of 3D data



Robotics



Augmented Reality



Autonomous driving



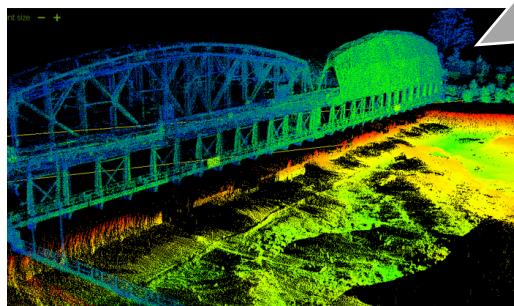
Broad Applications of 3D data



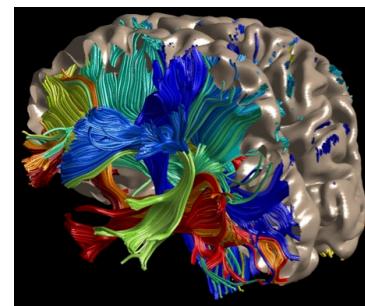
Robotics



Augmented Reality

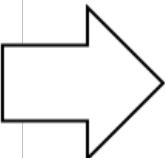


Autonomous driving



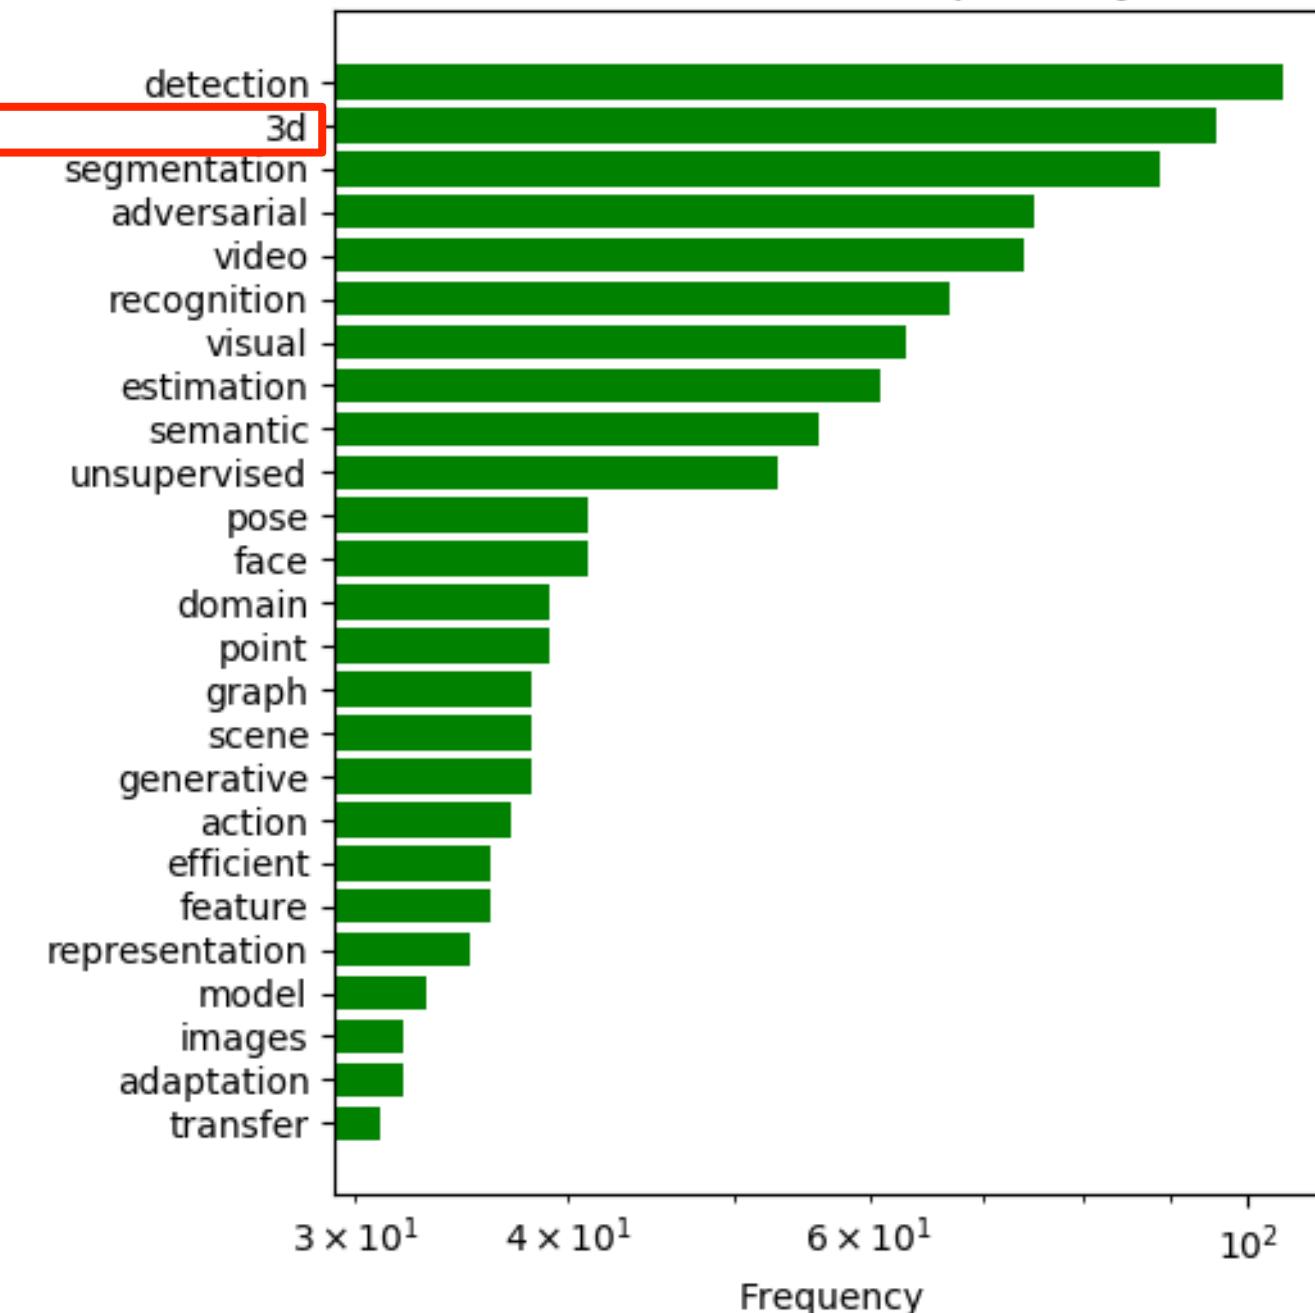
Medical Image Processing

Acquire Knowledge of 3D World by Learning



A priori knowledge of
the 3D world

CVPR 2019 Submission Top 25 Keywords

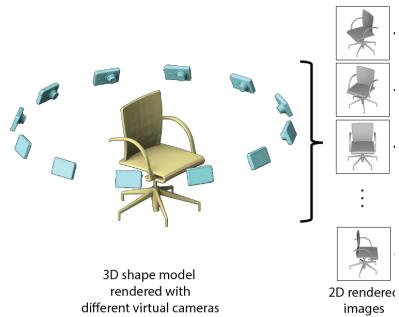


The Representation Challenge of 3D Deep Learning

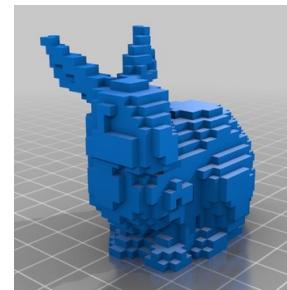
**Rasterized form
(regular grids)**

**Geometric form
(irregular)**

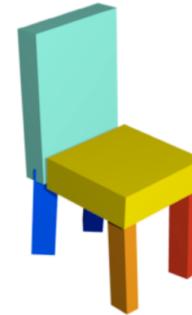
The Representation Challenge of 3D Deep Learning



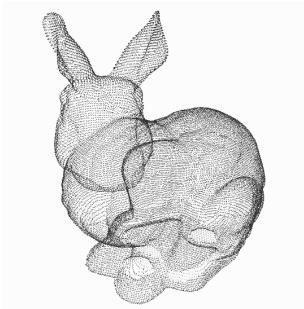
Multi-view



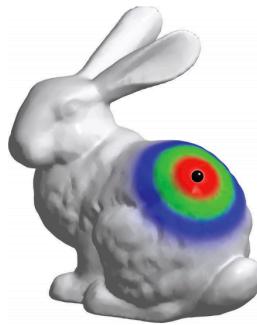
Volumetric



Part Assembly



Point Cloud



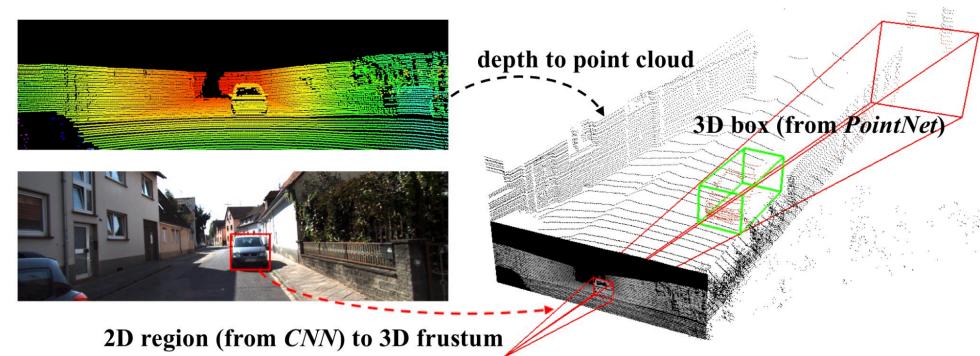
Mesh (Graph CNN)

$$F(x) = 0$$

Implicit Shape

The Richness of 3D Learning Tasks

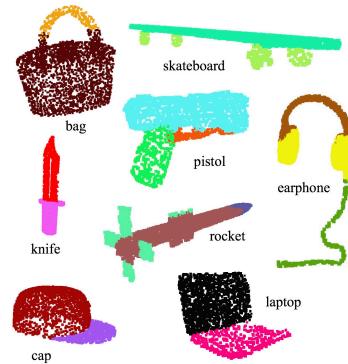
3D Analysis



Detection



Classification



Segmentation (object/scene)



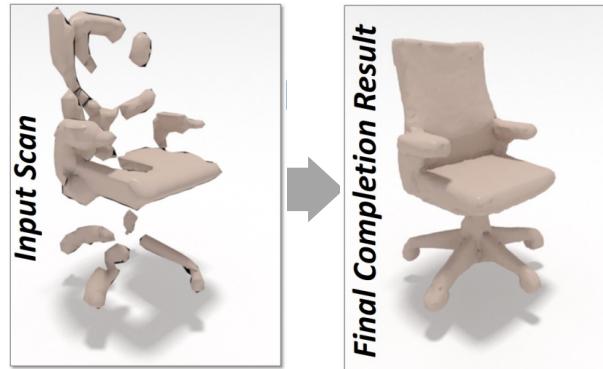
Correspondence

The Richness of 3D Learning Tasks

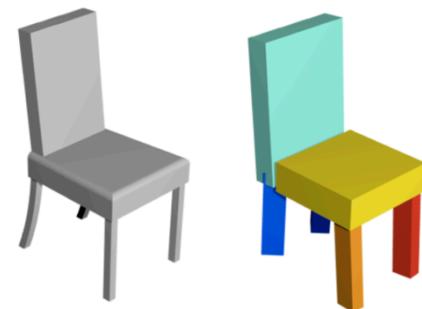
3D Synthesis



Monocular
3D reconstruction



Shape completion



Shape modeling

Agenda

- 3D Classification
- 3D Reconstruction

Volumetric CNN

Can we use CNNs but avoid projecting the 3D data to views first?

Straight-forward idea: Extend 2D grids 3D grids

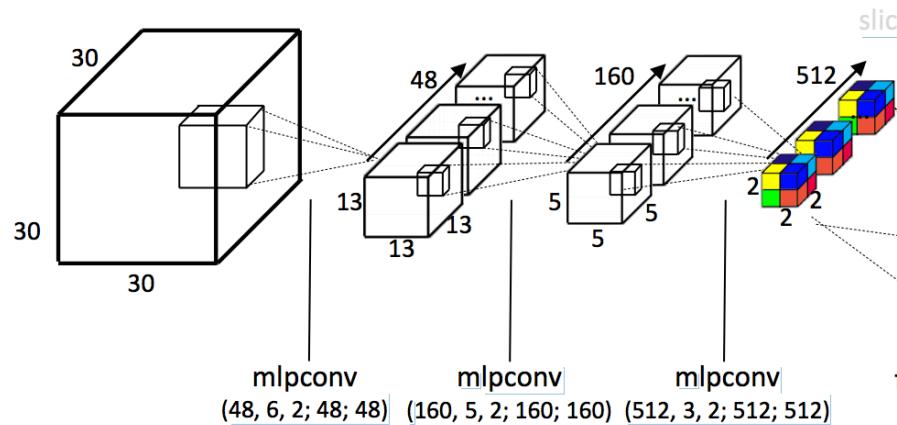
Voxelization

Represent the occupancy of regular 3D grids

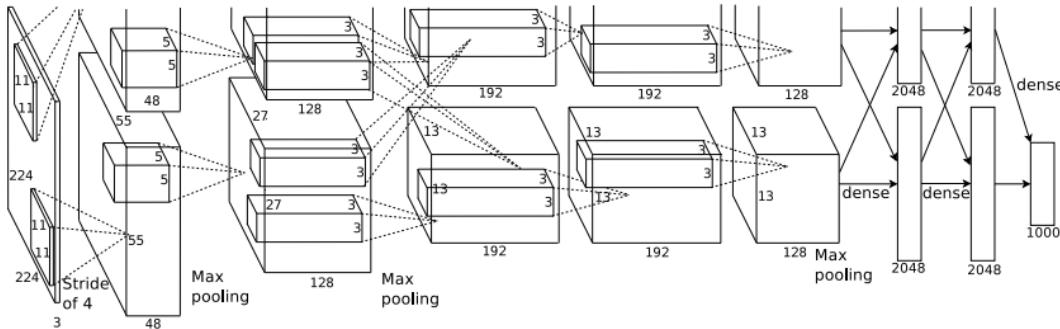


3D CNN on Volumetric Data

3D convolution uses 4D kernels



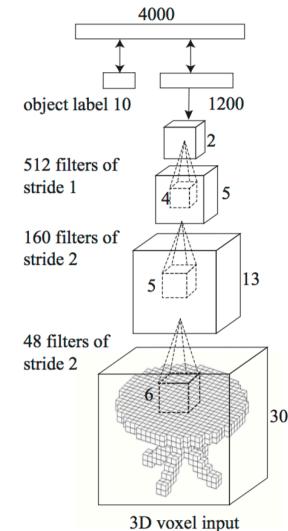
Complexity Issue



AlexNet, 2012

Input resolution: 224x224

$$224 \times 224 = 50176$$

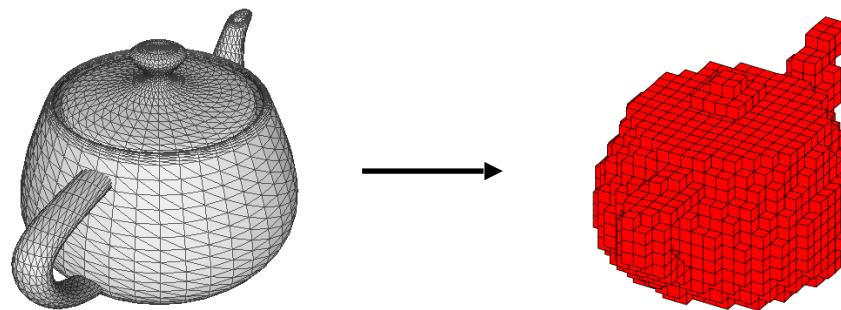


3DShapeNets,
2015

Input resolution: 30x30x30

$$224 \times 224 = 27000$$

Complexity Issue



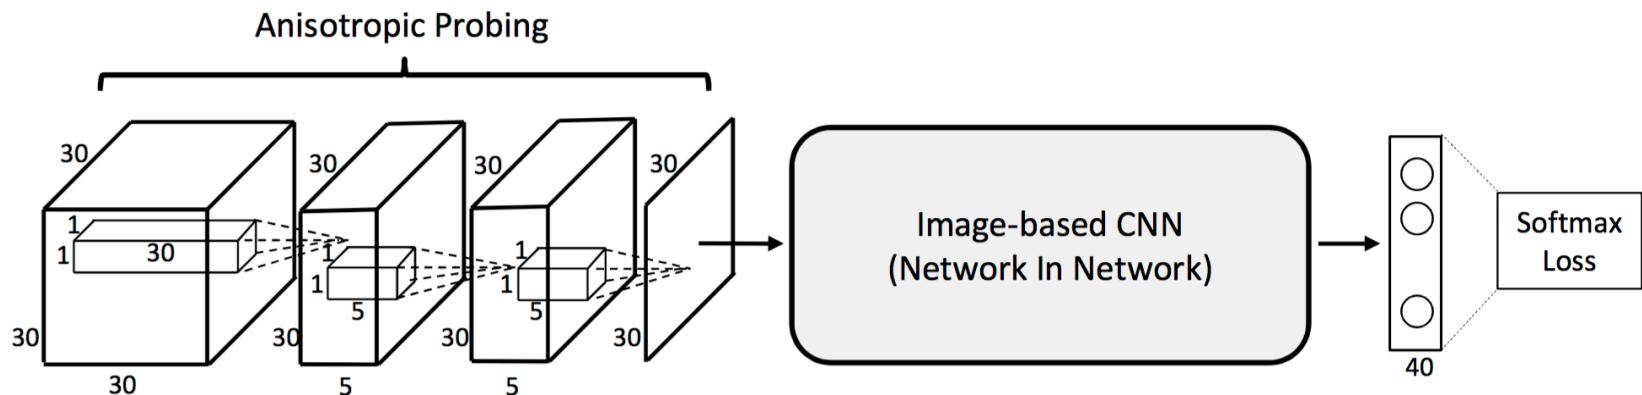
Polygon Mesh

Occupancy Grid
 $30 \times 30 \times 30$

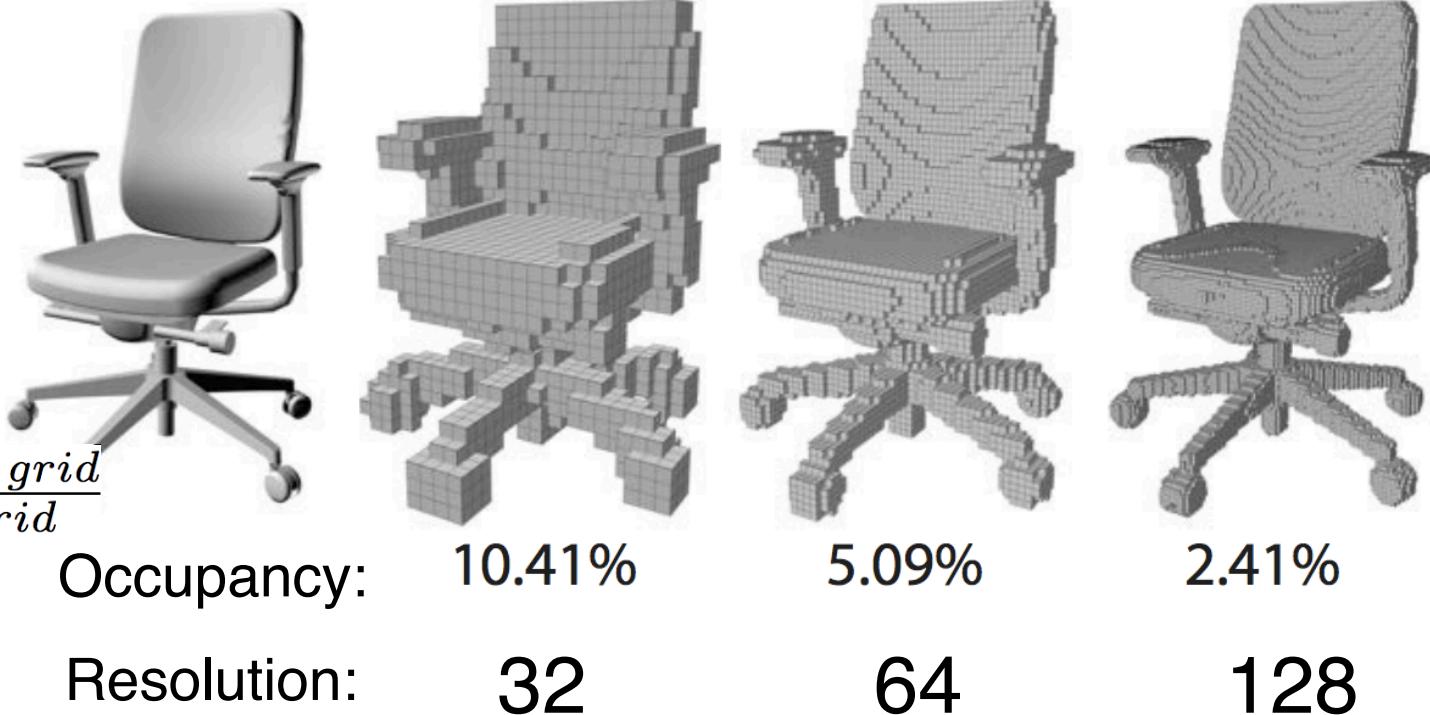
Information loss in voxelization

Idea 1: Learn to Project

*Idea: “X-ray” rendering + Image (2D) CNNs
very low #param, very low computation*

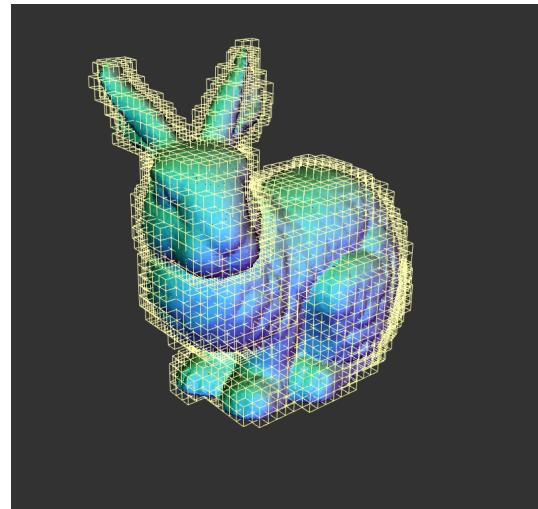
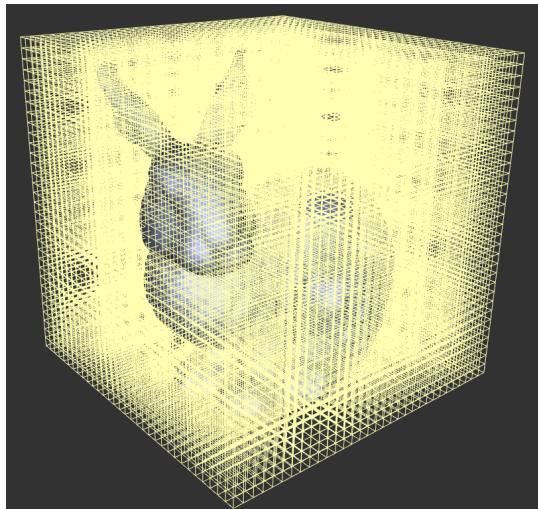


More Principled: Sparsity of 3D Shapes



Store only the Occupied Grids

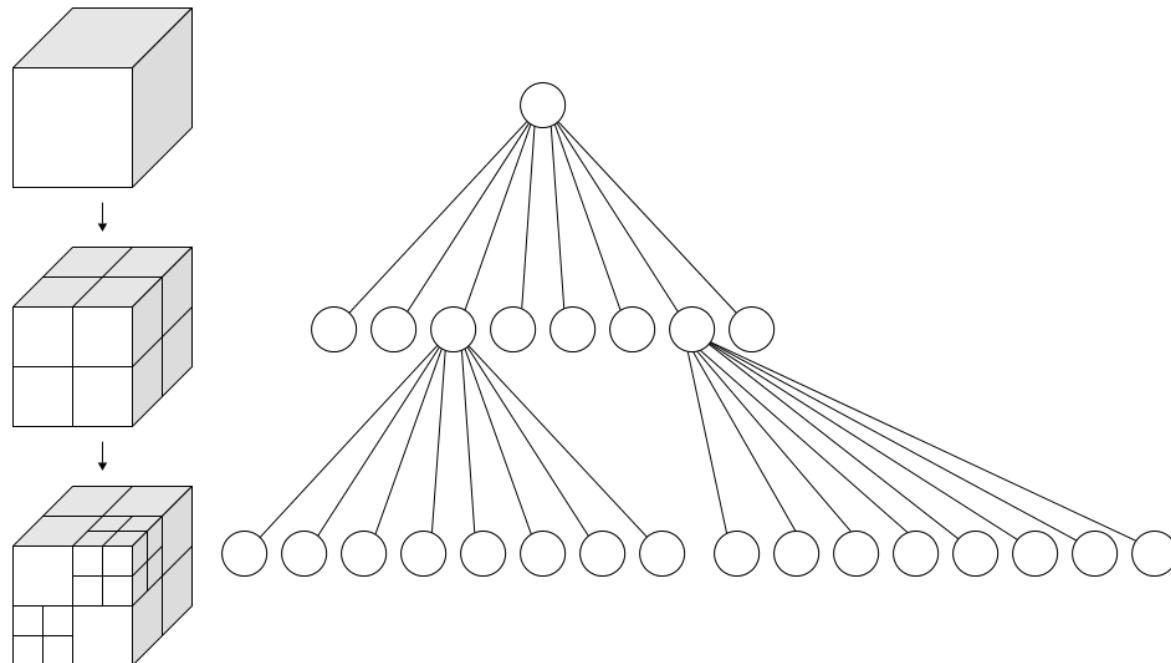
- Store the sparse surface signals
- Constrain the computation near the surface



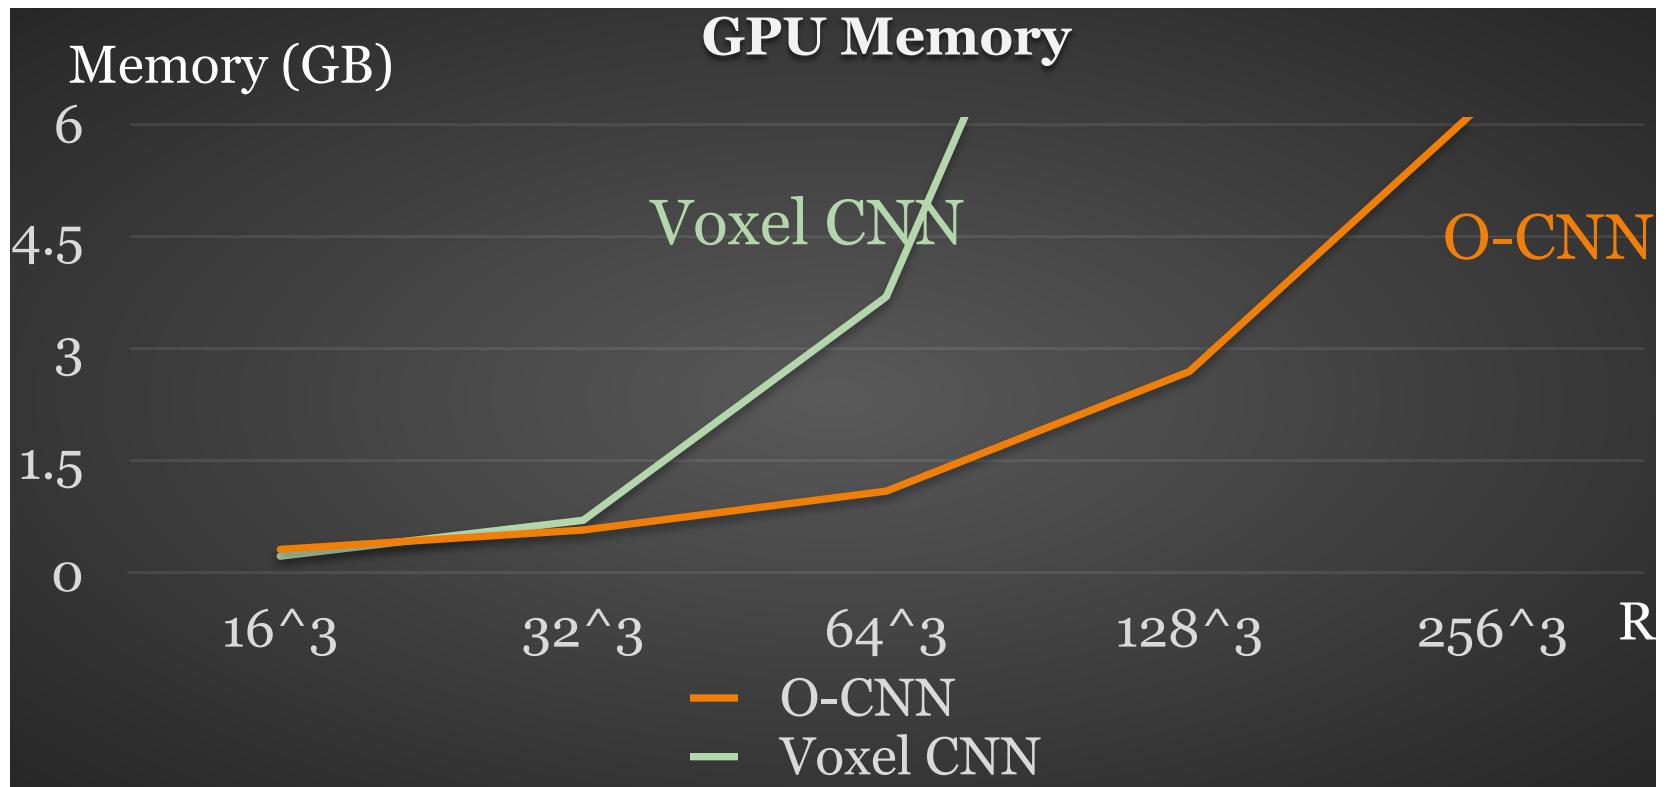
Octree: Recursively Partition the Space

Each internal node has exactly eight children

Neighborhood searching: Hash table



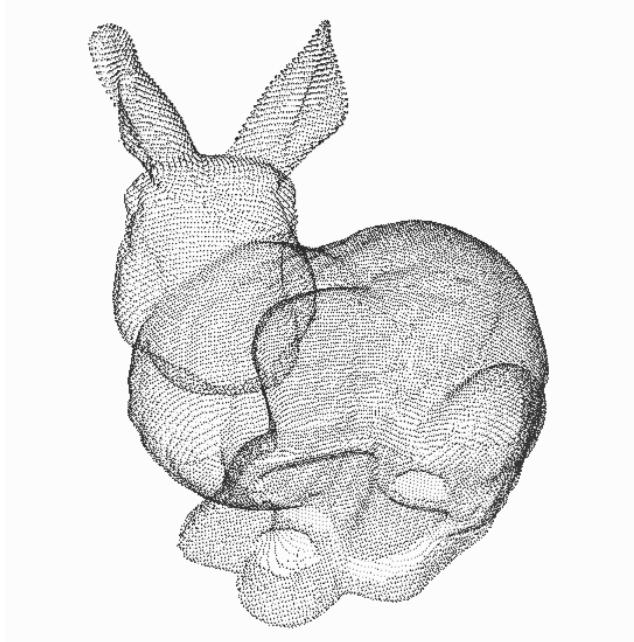
Memory Efficiency



Implementation

- SparseConvNet
 - [https://github.com/facebookresearch/
SparseConvNet](https://github.com/facebookresearch/SparseConvNet)
 - Uses ResNet architecture
 - State-of-the-art for 3D analysis
 - Takes time to train

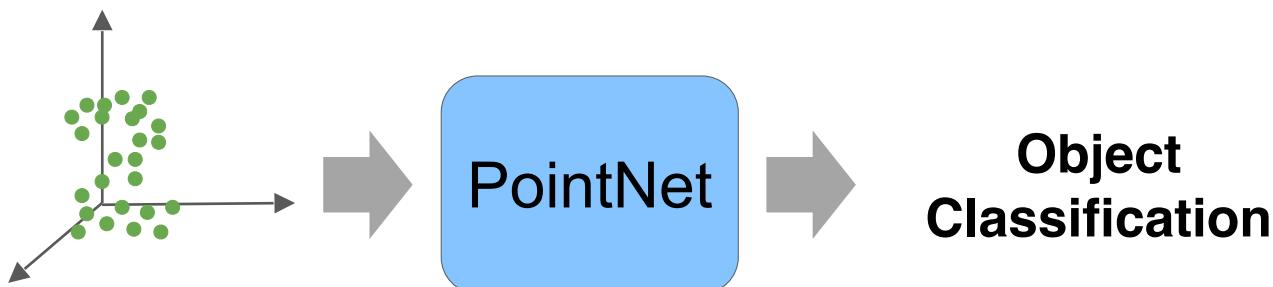
Point Networks



Point cloud
(The most common 3D sensor data)

Directly Process Point Cloud Data

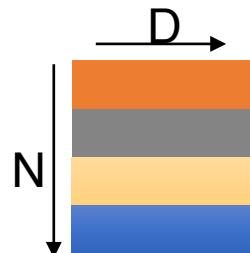
End-to-end learning for **unstructured**,
unordered point data



Qi, Charles R., et al. "Pointnet: Deep learning on point sets for 3d classification and segmentation", CVPR 2017
Zaheer, Manzil, et al. "Deep sets", NeurIPS 2017

Properties of a Desired Point Network

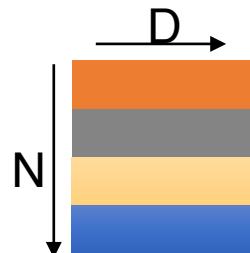
Point cloud: N **orderless** points, each represented by a D dim coordinate



2D array representation

Properties of a Desired Point Network

Point cloud: N **orderless** points, each represented by a D dim coordinate



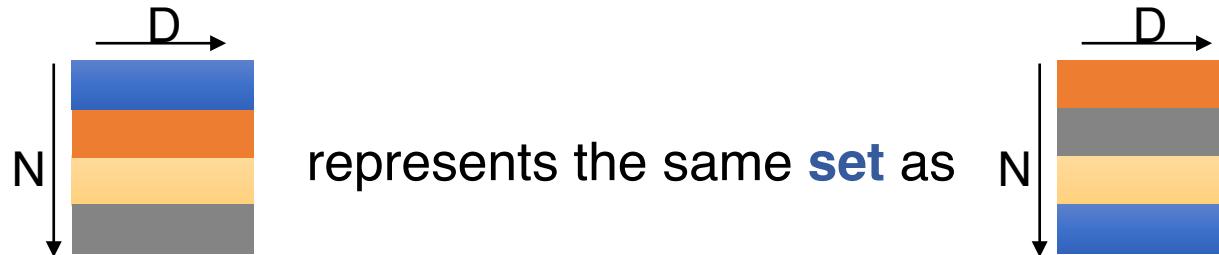
2D array representation

Permutation invariance

Transformation invariance

Permutation Invariance

Point cloud: N **orderless** points, each represented by a D dim coordinate



2D array representation

Permutation Invariance: Symmetric Function

$$f(x_1, x_2, \dots, x_n) \equiv f(x_{\pi_1}, x_{\pi_2}, \dots, x_{\pi_n}), \quad x_i \in \mathbb{R}^D$$

Examples:

$$f(x_1, x_2, \dots, x_n) = \max\{x_1, x_2, \dots, x_n\}$$

$$f(x_1, x_2, \dots, x_n) = x_1 + x_2 + \dots + x_n$$

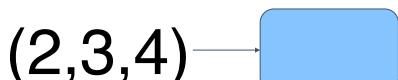
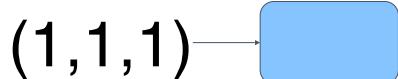
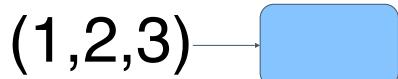
...

Construct a Symmetric Function

Observe:

$f(x_1, x_2, \dots, x_n) = \gamma \circ g(h(x_1), \dots, h(x_n))$ is symmetric if g is symmetric

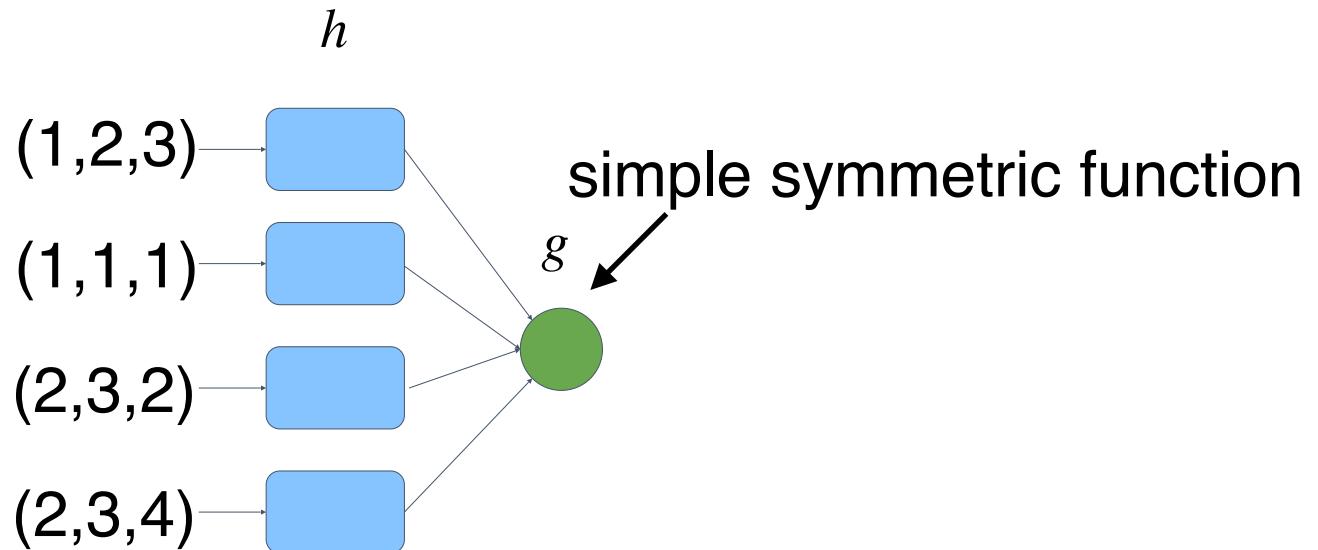
h



Construct a Symmetric Function

Observe:

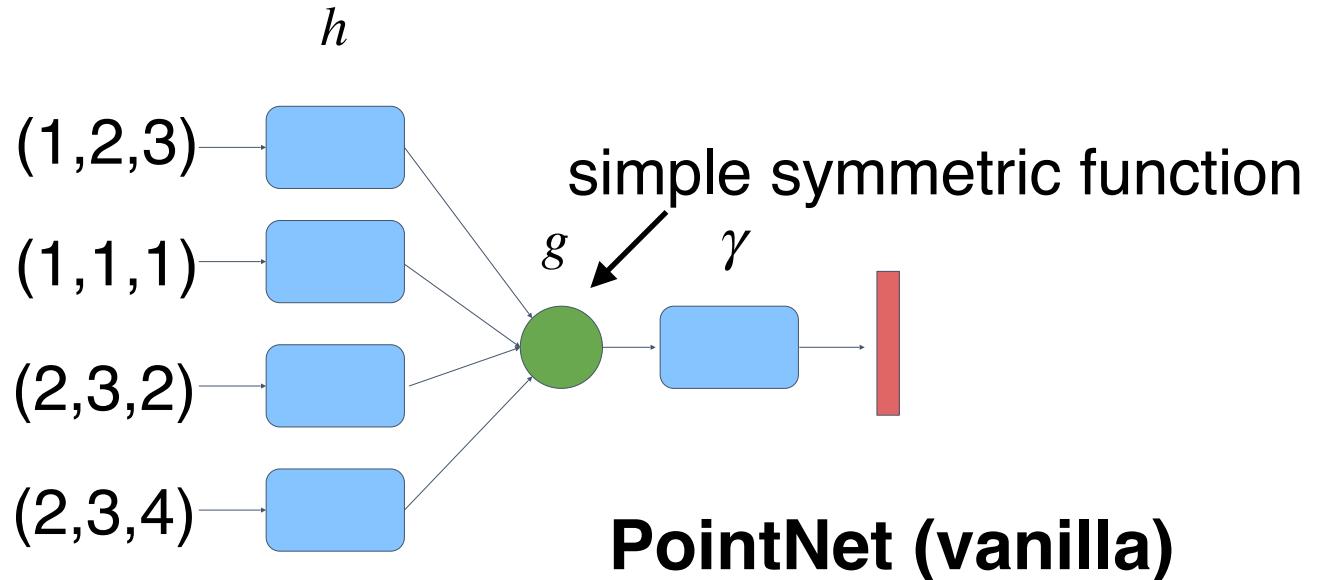
$f(x_1, x_2, \dots, x_n) = \gamma \circ g(h(x_1), \dots, h(x_n))$ is symmetric if g is symmetric



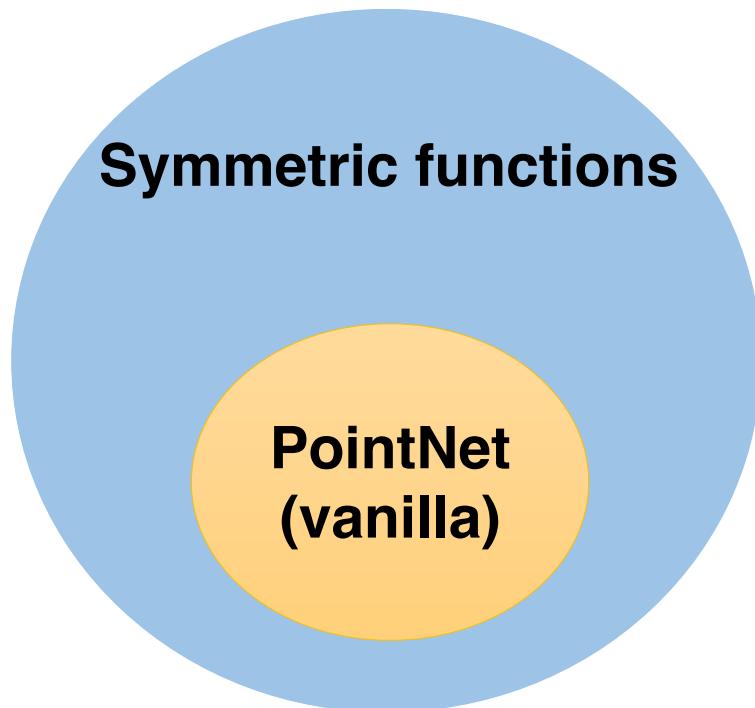
Construct a Symmetric Function

Observe:

$f(x_1, x_2, \dots, x_n) = \gamma \circ g(h(x_1), \dots, h(x_n))$ is symmetric if g is symmetric



Q: What Symmetric Functions Can Be Constructed by PointNet?

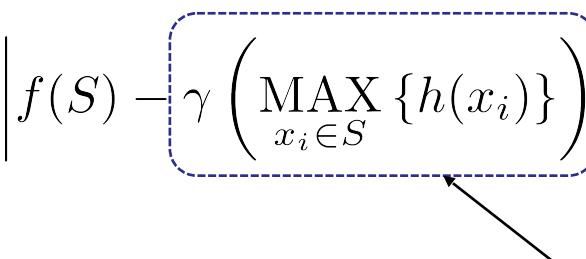


Universal Approximation Theorem

- Can approximate any “continuous” functions over sets
- “Continuous”: A function value would change by little if the point positions vary by little

$$\left| f(S) - \gamma \left(\text{MAX}_{x_i \in S} \{h(x_i)\} \right) \right| < \epsilon$$

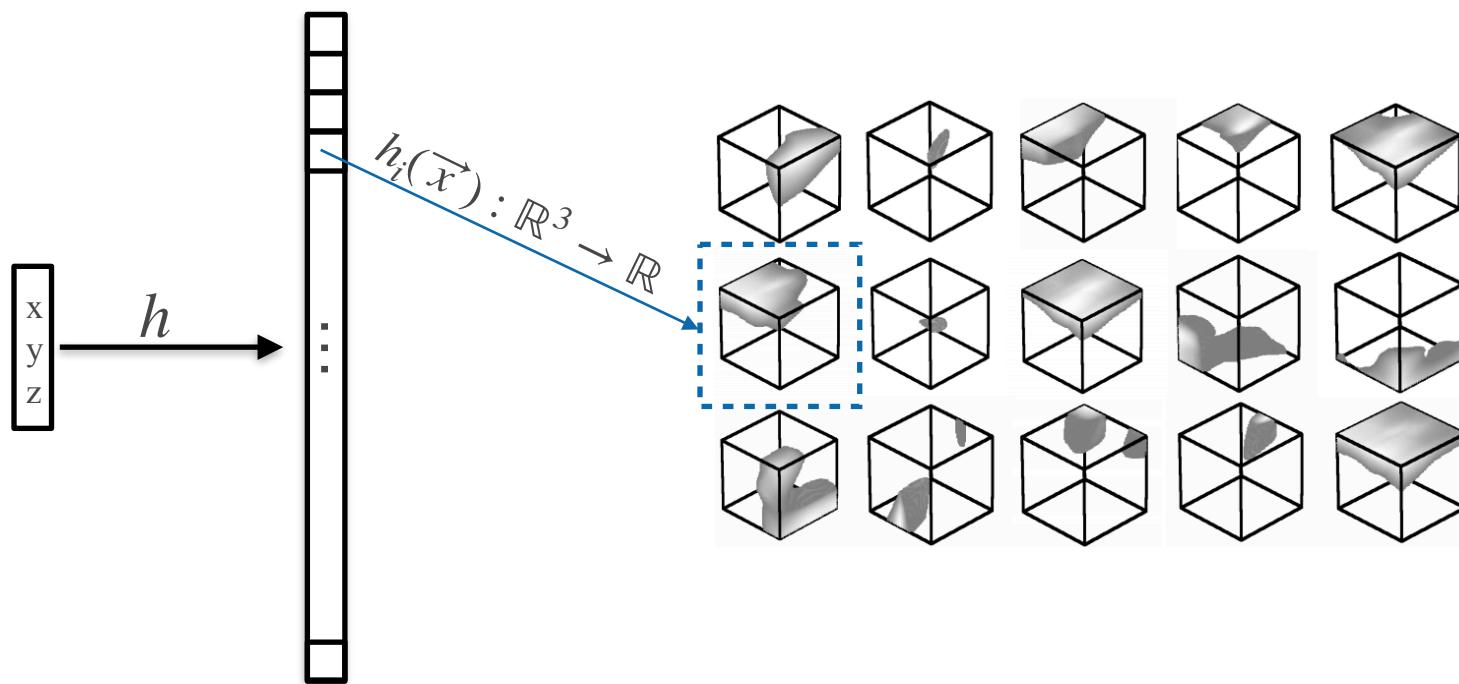
$S \subseteq \mathbb{R}^d,$ **PointNet (vanilla)**



Interpretation to “First Layer” Output

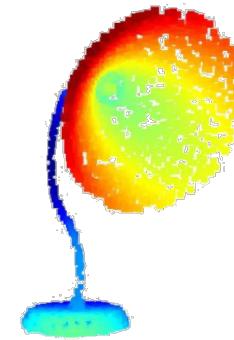
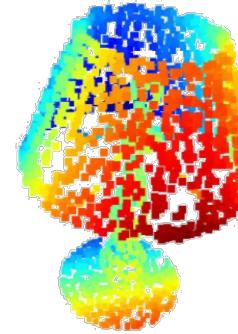
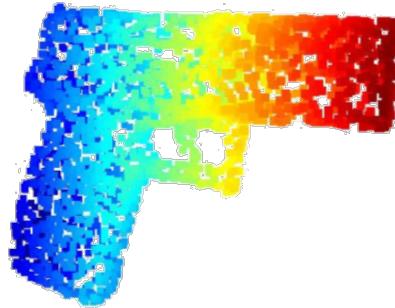
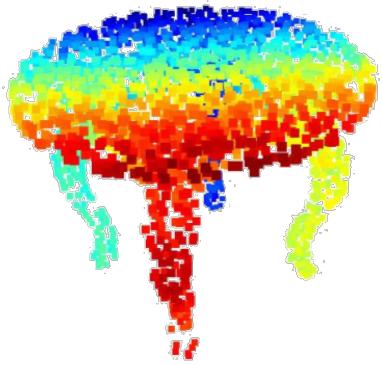
- Think of each dimension as a “binary” variable (the truth is a soft version)
- It encodes whether the point is in a certain spatial region
- The shape of the spatial region is learned

3D voxels of irregular boundaries!

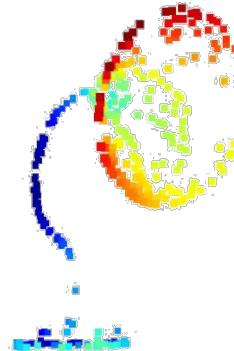
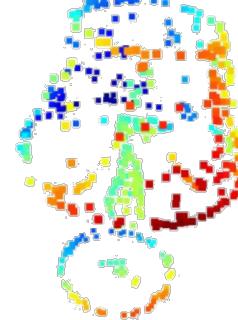
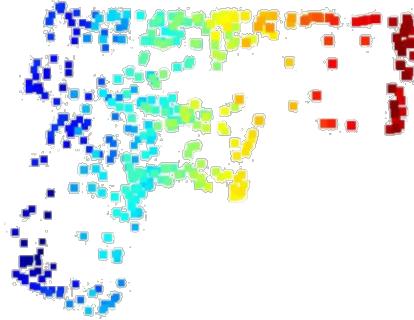
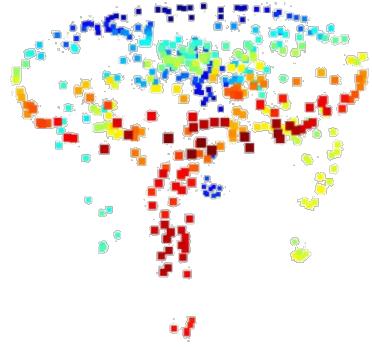


Salient Points: Points with Non-Zero Gradient w.r.t. Positions

Original Shape



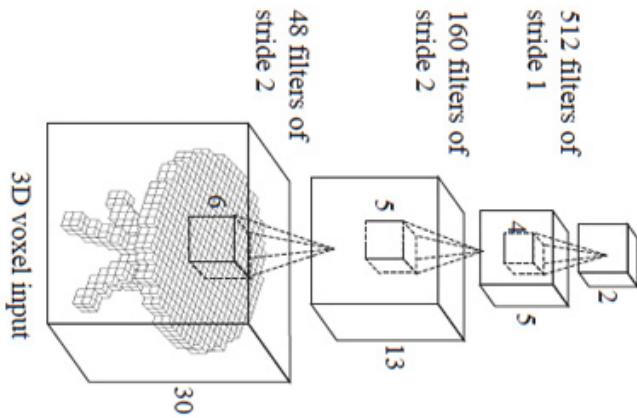
Critical Point Sets



Salient points are discovered!

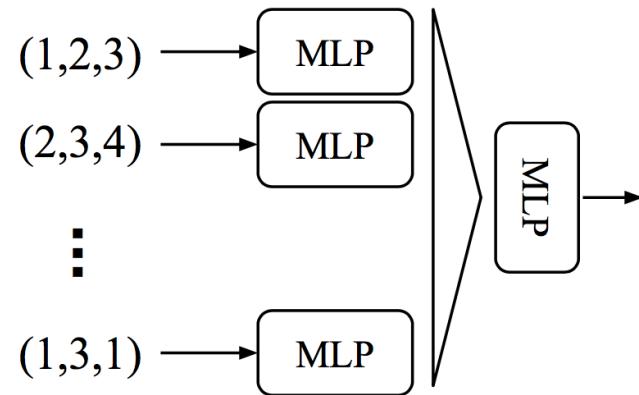
Limitations of PointNet

Hierarchical feature learning
Multiple levels of abstraction



3D CNN (Wu et al.)

Global feature learning
Either one point or all points



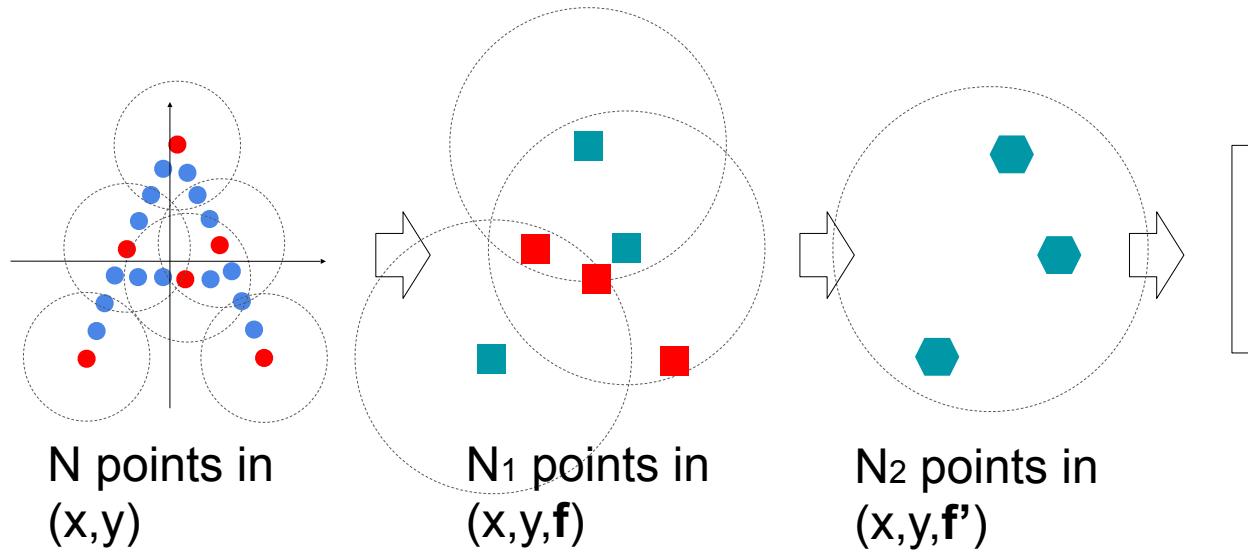
PointNet (vanilla) (Qi et al.)

- No local context for each point!
- Global feature depends on absolute coordinate. Hard to generalize to unseen scene configurations!

Points in Metric Space

- Learn “kernels” in 3D space and conduct convolution
- Kernels have compact spatial support
- For convolution, we need to find neighboring points
- Possible strategies for range query
 - Ball query (results in more stable features)
 - k-NN query (faster)

PointNet v2.0: Multi-Scale PointNet

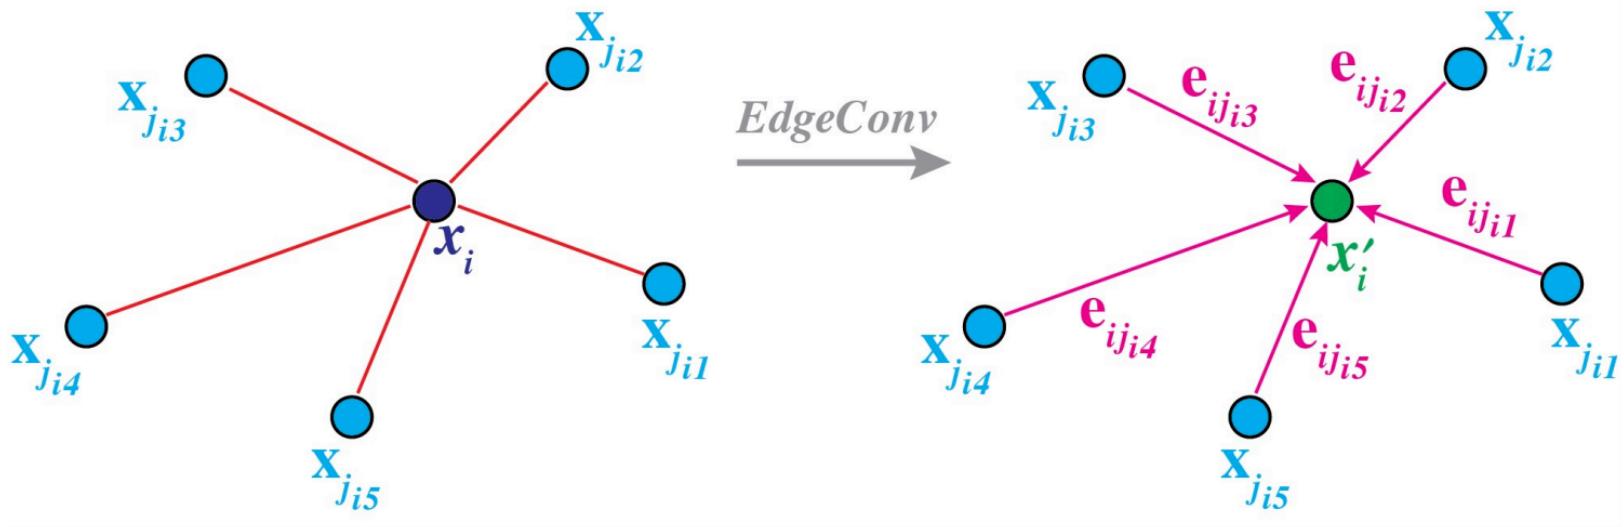


Repeat

- Sample anchor points
- Find neighborhood of anchor points
- Apply PointNet in each neighborhood to mimic convolution

Point Convolution As Graph Convolution

- Points -> Nodes
- Neighborhood -> Edges
- Graph CNN for point cloud processing



Wang et al., “Dynamic Graph CNN for Learning on Point Clouds”,
Transactions on Graphics, 2019

Liu et al., “Relation-Shape Convolutional Neural Network for Point Cloud Analysis”, *CVPR 2019*

Agenda

- 3D Classification
- 3D Reconstruction

From Single Image to Point Cloud

- It is possible to generate a **set** (permutation invariant)

